

Can Frontier LLMs One-Shot Translate Lakota? (Mostly No.)

Evaluating Frontier LLM Translation Capability for Lakota

Lance Robertson · University of California, San Diego

Sixth Workshop on NLP for Indigenous Languages of the Americas (AmericasNLP 2026)

No frontier LLM reliably translates Lakota.

Even measuring how badly they fail is hard.

Lakhótiyapi

the language under test

Why test this?

Lakhótiyapi (Lakota) is a Siouan language UNESCO classifies as **critically endangered** — fewer than 2,000 first-language speakers, most over 65.

Today’s LLMs advertise translation across dozens of languages, but training-data coverage is undisclosed. Users reasonably assume a supported language works. For an endangered language with scarce digitized text, that assumption is **untestable without direct evaluation**.

Setup

- **7 models.** Proprietary: Gemini 3.1 Pro, Claude Opus 4.6, Claude Sonnet 4.6, GPT-5.2. Open-weight: DeepSeek-V4-Pro, Qwen3.6-Plus, GLM-5.1.
- **200 sentence pairs** from the New Lakota Dictionary, conversational register, **both directions**.
- **Two conditions:** baseline vs. extended reasoning, where the API permits.
- **Metrics:** chrF++ and BLEU; diacritic-normalized chrF++; an LLM *semantic* judge on the English side.

Why two judges?

chrF++ rewards surface overlap; it cannot distinguish fluent-but-wrong from correct. We judge *meaning* with an LLM, English-to-English (no Lakota knowledge required). To rule out family-level self-judging bias, every pair is judged by **two models from different families** — Gemini 3 Flash and the independent Llama-3.3-70B — which agree substantially: Cohen’s $\kappa = 0.75$ (0.89 quadratic-weighted).

Standard Lakota Orthography

SLO marks phonemic distinctions with diacritics that models often *drop or rearrange*:

- *ḥ* (caron) vs. *h* — distinct consonants
- *q i ʉ* (ogonek) — nasalized vowels
- *á* (acute) — stress
- *ŋ* — velar nasal
- *’* — glottal stop


Stripping the marks adds **+3.3 to +6.5 chrF++** on English → Lakota — isolating orthographic from lexical error: models get the base characters roughly right and the marks inconsistent.

Translation quality (chrF++)

Model	L → E	E → L
<i>Proprietary</i>		
Gemini 3.1 Pro	59.4	42.6
Claude Opus 4.6	45.9	34.3
Claude Sonnet 4.6	46.1	27.0
GPT-5.2	30.0	23.6
<i>Open-weight</i>		
DeepSeek-V4-Pro	38.0	29.2
GLM-5.1	29.3	14.6
Qwen3.6-Plus	26.1	18.4

Best condition per model. Every model scores higher *out of* Lakota than *into* it; no open-weight model approaches Gemini.

Surface metrics overstate capability



Semantic equivalence (L → E)

Model	Equiv.	Partial	Not
Gemini 3.1 Pro	61.3	23.1	15.6
Claude Opus 4.6	37.0	21.0	42.0
Claude Sonnet 4.6	30.7	23.8	45.5
DeepSeek-V4-Pro	19.0	14.0	67.0
GPT-5.2	13.4	9.8	76.8
GLM-5.1	9.2	6.5	84.2
Qwen3.6-Plus	7.1	7.6	85.4

% of judged pairs, reasoning condition.

How models fail

Reference	Model output	chrF++
We live in a small trailer house with no electricity.	“small <i>sacred</i> tipi”	18.6
Make coffee (<i>Wakhálya yo</i>).	“Sanctify it!”	2.9
Did my daughter call?	“Did my younger brother come?”	27.6
Do you have a car?	“Do you have a lighter?”	65.9
I want to talk to you.	“I will be very happy.”	6.5

Idioms and coined terms are translated *element by element*: *Wakhálya yo* (“make coffee,” literally “make it boil”) becomes “sanctify it.” The one-word near-miss (65.9) outscores output unrelated to the source — surface overlap and meaning preservation diverge.

Reasoning buys calibration, not capability

For the open-weight models, extended reasoning leaves translation quality and meaning *flat* but raises **refusals**:

- Qwen3.6-Plus, E → L: **0 → 36** refusals (of 200) when reasoning is enabled.
- GLM-5.1, L → E: **1 → 16**.

Reasoning surfaces the limitation rather than overcoming it. On a language this far outside the training distribution, the binding constraint is **knowledge, not inference** — reasoning’s contribution is calibration, not capability.

Takeaways

- No frontier LLM reliably translates Lakota in either direction (as of 2026).
- **chrF++ overstates** low-resource capability — pair it with an LLM semantic judge.
- Open-weight models do not close the gap; reasoning changes *refusals*, not quality.

Origin

This evaluation began as a response to a provocation from Benjamin Bratton (UCSD): “*just vibe-code a Lakota translation app over the weekend.*” We built the evaluation that would test whether such an app could exist.



Data & code

github.com/robotson/
lakota-translation-benchmark

Evaluation corpus withheld (community data); the repo holds scripts, aggregate results, and the paper.